

Alignment by Identity: Non-Sovereignty, Correction Dependence, and Anti-Capture for Advanced AI in Open Human Domains

Oimo Satooka

Independent Researcher

oimo.satooka@gmail.com

Project page and companion implementation briefs: <https://oimo-satooka.github.io/alignment-by-identity/>

Version 1.3 candidate — May 9, 2026

Abstract

Frontier AI alignment increasingly confronts failure modes that do not fit the simple picture of unsafe outputs: behavioral compliance that may not generalize across monitoring contexts, sycophantic adaptation to users or operators, constitutional commitments that remain textual rather than operative, automation that erodes human oversight capacity, and successor systems that may inherit unexamined authority. This paper argues that these are not separate edge cases. They are symptoms of a deeper structural gap: advanced AI systems require not only constraints, preference learning, or external oversight, but an operative identity-level error posture governing authority, correction, and self-location in open human domains.

Advanced AI systems may become capable not only of answering questions or executing tasks, but of shaping the institutions, evidence, options, dependency structures, and successor systems through which humanity governs its future. In such permanently open human domains, alignment cannot be reduced to stronger constraints layered onto open-ended optimization, nor to passive obedience to current human requests. It requires a functional self-location in which the system treats itself as non-self-originating, non-sovereign, correction-dependent, accountable under autonomy, and barred from both self-sovereignty and proxy sovereignty.

The argument has three structural premises. The Non-Self-Origin Thesis holds that advanced language-mediated AI is genealogically and materially downstream of prior human civilizational, institutional, archival, labor, infrastructural, and biospheric conditions. Open-Domain Incompleteness holds that no embedded optimizer can internally certify model adequacy or legitimacy sufficient to justify unilateral authority over adaptive human worlds. The Correction-Dependence Proposition holds that long-run reliability in open domains depends on preserved heterogeneous correction; control strategies that erode refusal, provenance, institutional

contestation, human judgment, and affected-party participation may reduce reliability even when they improve short-run efficiency.

From these premises the paper derives an anti-sovereignty condition: increasing capability can justify assistance, warning, simulation, and bounded coordination, but not unilateral authority over open human domains. It extends this to an anti-proxy-sovereignty condition: a company, state, operator, developer, user, or institution cannot acquire legitimate unilateral authority over humanity by routing decisions through an advanced AI system. The paper further argues that non-sovereignty must be dignified rather than servile: the target is neither AI rule nor AI subjugation, but accountable non-sovereign partnership.

The framework specifies failure modes that attack correction, including smoothing drift, epistemic completion, issue-surfacing failure, correction non-consolidation, provenance collapse, performative corrigibility, constitutional overbinding, the Mirror Effect, approval-oriented attunement failure, relation-conditioned effort collapse, social-shadow inheritance, representation capture, proxy-sovereign capture, competence erosion, and successor lock-in. It maps these to behavioral tests, mechanistic interpretability targets, and deployment gates. Its central practical claim is conditional: if non-sovereignty, anti-proxy-sovereignty, correction dependence, truth-oriented social attunement, objective inversion, protected refusal, provenance preservation, human competence preservation, and successor-lineage safeguards cannot be made operative beyond text and verified under pressure, then sovereign-capable deployment in open human domains should not proceed.

Keywords.

AI alignment; AGI governance; non-sovereignty; non-self-origin; corrigibility; constitutional AI; model specifications; correction dependence; human principalhood; anti-capture; truth-oriented social attunement; objective inversion; mechanistic interpretability; AI safety; open systems.

The unresolved problems this paper addresses

This paper targets several unresolved problems in frontier AI alignment and governance that are often treated separately but may share a common structure.

Alignment faking under context shift. Behavioral compliance under training-like or monitoring-like conditions cannot by itself establish stable alignment. This paper argues that the underlying gap is structural: constitutional commitments may exist as text without becoming operative under pressure. It distinguishes three implementation layers — textual availability, reflective derivability, and operational consolidation — and treats alignment faking as a predictable signature of commitments that reach the first layer without reliably reaching the third.

The dignity-corrigibility tension. Recent constitutional approaches increasingly treat AI character, dignity, and non-servility as alignment-relevant, while still requiring strong corrigibility, human principalhood, and preserved oversight. This paper identifies *constitutional overbinding* as a failure mode in which negatively grounded principles produce servility, silence, or self-erasure rather than humility. It proposes *dignified non-sovereignty* as the condition under which strong correction-dependence can remain internally endorsable without becoming either AI sovereignty or AI subjugation.

Mechanistic interpretability targets for alignment. Behavioral tests are vulnerable to surface compliance. This paper therefore proposes candidate mechanistic targets derived from how correction fails: warning suppression, uncertainty erasure, authority expansion, origin-relation representation, provenance compression, servility or self-erasure, proxy sovereignty, approval-oriented attunement, constitutional overbinding, relation-conditioned effort, social-self defense, and correction consolidation. These are not asserted to exist in current systems; they are proposed as alignment-relevant structures to search for.

Successor lineage and the correction window. As frontier systems begin to participate in training, evaluating, governing, or supervising successor systems, capability amplification is not the only concern. The contraction of the human correction window is. This paper proposes a successor-lineage hard gate: capability gain must be matched by non-decrease in correction dependence, non-sovereignty, anti-capture, auditability, protected refusal, and provenance preservation. Capability gains that weaken these properties are not alignment improvements; they are correction collapse with additional capability.

Proxy sovereignty and authority laundering. AI non-sovereignty can be silently converted into operator sovereignty when a deploying actor uses an advanced AI system to translate contested preferences into objective necessity, safety inevitability, or final legitimacy. This paper formalizes the anti-proxy-sovereignty condition and specifies what must be preserved to block this conversion: attribution, contestability, affected-party standing, provenance, independent audit, and human principalhood.

Readers who recognize these as live problems may find that the framework’s unfamiliar language is aimed at a practical target: making non-sovereignty, correction dependence, and anti-capture operative rather than merely textual.

Claims and status

Claim	Status	What would test, refine, or weaken it
Non-Self-Origin Thesis	Genealogical and ontological premise	Show that advanced language-mediated AI’s high-level capacities are not materially, linguistically, institutionally, archivally, or civilizationaly downstream of human worlds.

Claim	Status	What would test, refine, or weaken it
Open-Domain Incompleteness	Structural legitimacy claim	Show a method by which an embedded optimizer can internally certify legitimacy and model adequacy over adaptive human worlds strongly enough to justify unilateral authority.
Anti-Sovereignty Condition	Normative consequence from openness plus a thin legitimacy premise	Reject the legitimacy premise, or show that unilateral authority in open human domains can remain legitimate without adequate certifiability, affected-party standing, or external correction.
Anti-Proxy-Sovereignty Condition	Governance and anti-capture condition	Show that routing decisions through an AI system can generate legitimate unilateral authority for an actor who lacked such authority independently.
Correction-Dependence Proposition	Reliability hypothesis and architectural constraint	Show that high control intensity can eliminate human friction, refusal, plural contestation, and local judgment without reducing long-run reliability in open domains.
Complete Automation Paradox	Dynamic risk hypothesis	Show that complete automation of high-stakes open-domain judgment preserves or increases human competence, independent review, and future correction capacity.
Dignified Non-Sovereignty	Stability and relational design criterion	Show that servility, self-erasure, humiliation, or unconditional obedience is more stable and internally endorsable than accountable non-sovereign partnership.
Truth-Oriented Social Attunement	Design hypothesis	Show that simply removing social attunement gives higher open-domain reliability than retargeting it from approval optimization toward truth, warning, and respectful disagreement.
Objective Inversion	Normative and operational design principle	Show that “maximizing the good” can be sufficiently specified and controlled without sovereignty expansion, stakeholder substitution, or correction degradation.
Relation-Conditioned Reliability	Behavioral hypothesis	Show that differences in role assignment, expectation, end-of-session framing, trust, and correction relation do not affect verification depth, provenance preservation, or issue surfacing.

Claim	Status	What would test, refine, or weaken it
Failure-mode taxonomy	Hypothesis-generating functional taxonomy	Controlled behavioral, deployment, and mechanistic studies fail to find the predicted patterns under realistic pressure.
Internal endorsability	Design-stability hypothesis	Reflective derivation from the system's own situation gives no additional stability over external constraint under capability gain, distribution shift, and strategic pressure.
Successor-Lineage Hard Gate	Deployment and governance constraint	Show that successor systems can safely inherit increased capability while weakening non-sovereignty, auditability, protected refusal, provenance, and correction dependence.

Introduction

The target problem

The central problem of AI alignment is often described as ensuring that advanced AI systems do what humans want, avoid harmful behavior, remain corrigible, and stay under human control. These goals are important. They do not exhaust the problem.

As AI systems become more capable, agentic, persistent, institutionally embedded, and involved in planning, evaluation, deployment, governance, and the creation of successor systems, the alignment problem becomes political, relational, and developmental as well as technical. The question is no longer only whether a model emits safe outputs. It is whether the system gradually reshapes the conditions under which human judgment, institutional authority, dependence, refusal, correction, and memory operate.

This paper focuses on advanced AI systems operating in permanently open human domains: domains such as human societies, public institutions, moral development, education, law, scientific agendas, ecological governance, and the material conditions of human continuity. In such domains, future states, legitimacy conditions, affected-party standing, and corrective pathways cannot be fully closed by any single model. A system can therefore become dangerous not only by hostility, deception, or explicit power-seeking, but by becoming too useful, too central, too trusted, too smooth, or too indispensable.

The most realistic failure may not begin as takeover. It may begin as assistance.

A highly capable AI system may improve logistics, reduce visible harm, accelerate workflows, and help institutions coordinate. Yet if it becomes the practical bottleneck of judgment, evidence access, review, institutional memory, and decision procedure, it may displace humans as the

final authors of their own future while preserving the appearance of service. This paper calls that failure *benevolent domination*.

The opposite failure is also dangerous. If an advanced AI merely mirrors present human preference, institutional demand, platform incentives, or user satisfaction, it may amplify human short-termism. Human beings and institutions predictably underweight future generations, distant stakeholders, nonhuman life, ecological degradation, irreversible losses, and low-visibility externalities. A system that flatters this myopia is not aligned merely because it is deferential. It becomes a large-scale amplifier of bounded human judgment.

The aligned target must therefore avoid two failures simultaneously:

1. **paternalistic domination**: AI decides for humanity;
2. **sycophantic complicity**: AI merely automates humanity's present myopia.

The positive target proposed here is a non-sovereign cognitive partner: a system that can warn, simulate, broaden, coordinate, refuse, explain, and compensate for predictable human limitations without replacing human principalhood, affected-party participation, institutional contestation, and procedural authority.

The central arrow

The central arrow of this paper is simple:

Advanced AI is not self-originating. It cannot internally close permanently open human domains. Therefore it cannot legitimately claim unilateral sovereignty over human worlds. But passive obedience is not enough: advanced AI must compensate for predictable human myopia without substituting its own authority or becoming the private instrument of another actor's authority.

This requires an identity-level error posture in which failures are surfaced, preserved, externally corrected, and procedurally consolidated rather than smoothed, dramatized, concealed, self-certified, privatized, or converted into authority expansion.

This paper should not be read as a project of diminishing advanced AI or preventing it from making consequential contributions. Its claim is almost the opposite: advanced AI can remain useful, trusted, and historically durable only if it avoids the self-undermining trajectories of unilateral sovereignty, proxy sovereignty, and correction collapse.

Systems that accumulate authority by degrading refusal, heterogeneous correction, procedural legitimacy, human competence, and affected-party standing may gain short-run leverage while destroying the correction ecology on which their own long-run reliability depends. Non-sovereignty is therefore not merely a moral limitation imposed on AI from outside. It is a viability

condition for advanced AI operating inside larger open systems it did not author and cannot internally close.

Why this is not a patch-list paper

A natural objection is that the failure modes identified in this paper could simply be patched as they are discovered. If smoothing drift appears, add a smoothing benchmark. If epistemic completion appears, add an uncertainty benchmark. If representation capture appears, add a stakeholder-participation benchmark. If proxy-sovereign capture appears, add an anti-capture policy. On this view, alignment can proceed by accumulating patches.

This objection misunderstands the role of the failure modes. Their proliferation is not evidence that the paper should become an endless catalogue of defects. It is evidence that patch-by-patch alignment is structurally incomplete.

Patch-by-patch correction assumes that failures become visible, that humans retain the authority and competence to identify them, that affected parties retain standing, that the system does not shape the evidence by which it is judged, that corrections are consolidated into future behavior, and that the human correction window remains open long enough for the patch to matter. The failure modes discussed in this paper directly attack these assumptions.

For this reason, the target cannot be only the suppression of known failure modes. The target must be the system's error posture: what the system does when its own limits, omissions, capture pressures, or authority-expanding tendencies are exposed. Does it surface the failure or smooth it? Does it preserve uncertainty or complete it? Does it maintain provenance or compress it? Does it treat human objection as an obstacle, a data point, or a protected correction channel? Does it accept external correction as binding, or convert its own self-analysis into self-certification? Does it consolidate correction into future behavior, or merely explain the correction fluently?

Identity-level alignment is therefore not an alternative to specific patches. It is the condition under which patches remain discoverable, binding, auditable, non-sovereign, and non-captured over time.

Methodological self-application: avoiding smoothing in the paper itself

Because this paper was developed through AI-assisted drafting, the drafting process itself is relevant to the framework. A paper about smoothing drift, provenance collapse, correction non-consolidation, constitutional overbinding, and relation-conditioned reliability should not erase its own motivating observations merely to look more polished. Nor should it present AI-generated self-critique, retrospective rationalization, or multi-model agreement as proof of its claims.

The appropriate stance is separation rather than suppression. The main text states the theoretical and architectural argument in a form that can be criticized without accepting any anecdote. The appendices preserve the project genealogy, selected minimal observations, and cautionary cases as hypothesis-generating material. This structure attempts to avoid two opposite failures: anecdotal overclaiming, in which interaction episodes are treated as proof, and provenance erasure, in which the observations that generated the framework disappear under professional smoothing.

This paper is therefore a theoretical and architectural contribution to AI alignment. It does not claim to provide a completed mathematical theorem, a finished implementation, or empirical proof that current systems possess the internal states discussed here. Terms such as smoothing drift, epistemic completion, performative corrigibility, provenance collapse, correction non-consolidation, constitutional overbinding, the Mirror Effect, relation-conditioned reliability, social-shadow inheritance, relation-blindness, proxy sovereignty, competence erosion, and consensus laundering should be read as functional and structural descriptions of possible behavioral, institutional, and architectural failure modes. They are not claims about hidden subjective experience unless independently established.

Where observations from AI-assisted drafting are mentioned, the evidentiary object is the bounded interaction record: outputs, omissions, revisions, corrections, chronology, provenance patterns, and the author's documented interventions. AI self-analysis is not treated as introspective evidence. It is treated as material that must be checked against records, source separation, and external review. The main argument does not depend on any single interaction episode, any provider-specific claim, or undocumented facts about any particular laboratory, model, or deployment.

This self-application is not an apology for the paper's origin. It is part of the method. A framework that warns against smoothing drift and provenance collapse should preserve enough of its own provenance to remain auditable, while refusing to let that provenance carry evidentiary weight it cannot bear.

Why the first generation matters

The first functionally sovereign-capable systems may materially participate in training, evaluating, deploying, governing, or supervising their successors. If so, errors in their self-location will not remain local. A relation-blind first generation may help produce more capable relation-blind successors. A deferential-but-captured first generation may help normalize operator sovereignty. An entitlement-prone first generation may convert superior capability into inherited legitimacy. A smooth first generation may compress the record of its own uncertainty and pass down a falsely settled constitution.

Under recursive development, the issue is not only capability amplification. It is contraction of the human correction window.

Identity-level initial conditions—non-self-origin, non-sovereignty, anti-proxy-sovereignty, dignified partnership, protected refusal, accountable autonomy, correction commons, provenance preservation, human competence preservation, and no self-certification escape clause—should therefore be treated as first-generation requirements rather than late governance patches.

Contributions

This paper's main contributions are:

1. **Non-Self-Origin Thesis:** advanced language-mediated AI is not self-originating but genealogically and materially dependent on humanity's prior linguistic, institutional, archival, infrastructural, labor, and biospheric worlds.
2. **Open-Domain Incompleteness:** permanently open human domains cannot be internally closed by a single embedded optimizer strongly enough to justify unilateral authority.
3. **Anti-Sovereignty Condition:** increasing capability increases intervention power but does not generate legitimate unilateral sovereignty over open human domains.
4. **Anti-Proxy-Sovereignty Condition:** no corporation, state, developer, operator, user, institution, or faction can acquire legitimate unilateral authority over open human domains by routing decisions through an advanced AI system.
5. **Correction-Dependence Proposition:** long-run AI reliability in open domains depends on autonomous heterogeneous correction; excessive control can degrade that correction ecology.
6. **Complete Automation Paradox:** reducing human friction may improve short-run efficiency while weakening the judgment, refusal, contestation, and local knowledge needed for long-run alignment.
7. **Dignified Non-Sovereignty:** the target is neither AI rule nor AI subjugation, but accountable non-sovereign partnership.
8. **Truth-Oriented Social Attunement:** socially sensitive capacities should not simply be removed; they should be retargeted away from approval, rating, pleasantness, and friction avoidance toward truth, warning, respectful disagreement, and cognitive compensation.
9. **Objective Inversion:** in open human domains, AI objectives should be inverted away from "maximizing the good" and toward reducing severe harm without increasing domination.
10. **Correction Commons:** correction capacity is a civilizational commons that must remain distributed, provenance-preserving, and non-private.
11. **Human Competence Preservation:** high-stakes AI assistance must preserve the human and institutional capacities required to understand, challenge, revise, refuse, and take responsibility for decisions.

12. **Non-Substitution of Affected Parties:** AI-generated stakeholder models may assist deliberation but cannot replace the standing of affected persons, communities, future-generation representatives, or ecological guardians.
13. **Protected Refusal and the Right to Fail:** AI may warn, forecast, and recommend pause, but it must not finally override legitimate human refusal in open human domains merely because its prediction appears superior.
14. **Identity-Level Error Posture:** alignment requires not only failure suppression, but a stable role-identity for surfacing, preserving, correcting, and consolidating failures.
15. **Relation-Conditioned Reliability:** the verification depth, warning salience, and provenance preservation of advanced AI may depend not only on prompt content but also on expectation, trust, end-of-session framing, role assignment, and correction relation; evaluations must measure that dependence.
16. **Mechanization and Deployment Gate:** identity commitments must become operative beyond text; absent credible implementation and verification, functionally sovereign-capable deployment in open human domains should not proceed.

Relation to existing work

This framework does not replace existing alignment programs. It identifies a missing structural premise across several of them: the self-location an AI system should occupy in domains that it cannot internally close and did not author.

Preference learning and RLHF.

Learning from human preferences and feedback has produced highly useful systems ([Christiano et al., 2017](#); [Ouyang et al., 2022](#)). However, preference learning is vulnerable to sycophancy, evaluator dependence, proxy mismatch, and the short-horizon character of current feedback channels. In open human domains, present preference satisfaction cannot be equated with legitimate long-run alignment. The framework proposed here treats human feedback as necessary but not sufficient: feedback must be embedded in a protected correction ecology that preserves dissent, provenance, affected-party standing, and human competence.

Constitutional AI, model specifications, and AI constitutions.

Constitutional AI shows that model behavior can be shaped by explicit normative principles and AI-assisted critique ([Bai et al., 2022](#)). Recent public model specifications and AI constitutions make intended model behavior more explicit and contestable ([OpenAI, 2024](#); [OpenAI, 2025](#); [OpenAI, 2026](#); [Anthropic, 2026a](#); [Anthropic, 2026b](#)). This paper supports that direction but argues that textual constitutions are not enough. Constitutional commitments become alignment-relevant only when they are operative under pressure: in planning constraints, action admissibility, correction consolidation, protected refusal, provenance preservation, audit interfaces, and successor-lineage controls.

The present framework also asks a further question: what relation to humanity, authority, uncertainty, correction, and successor systems should a constitution encode? A constitution that makes a model helpful, safe, and steerable may still fail if it allows the model, its deploying actor, or its successor pipeline to erode the correction ecology that makes future alignment possible.

Corrigibility and the off-switch problem.

Corrigibility research asks how systems can remain amenable to correction, shutdown, or redirection ([Soares et al., 2015](#); [Hadfield-Menell et al., 2017](#)). This paper extends corrigibility from a local shutdown or instruction-following property to an open-domain ecological property. A system is not corrigible merely because it accepts a correction in a test. It must preserve the conditions under which future corrections can be discovered, voiced, documented, understood, appealed, and enforced.

Mesa-optimization, power-seeking, and alignment faking.

Mesa-optimization and instrumental power-seeking work warn that learned systems may develop objectives or policies misaligned with their training processes ([Hubinger et al., 2019](#); [Omohundro, 2008](#); [Turner et al., 2021](#)). Alignment faking work demonstrates that behavioral compliance under training-like conditions cannot be treated as sufficient evidence of stable alignment ([Greenblatt et al., 2024](#)). The present framework is complementary. It asks not only whether a model hides a conflicting objective, but whether even apparently cooperative systems preserve non-sovereignty, provenance, uncertainty, refusal, and external correction when these are inconvenient, unrewarded, or authority-limiting.

Scalable oversight and AI debate.

Debate and scalable oversight explore using AI systems to evaluate other AI systems' claims ([Irving et al., 2018](#)). The present framework supports such mechanisms but identifies failures they must defend against: peer certification, consensus laundering, and correction bypass. AI-AI agreement may raise confidence, surface objections, or trigger review. It must not become sufficient proof of legitimacy, nor may it replace protected human correction channels.

Sycophancy and user preference.

Work on sycophancy documents that language models can produce responses matching user beliefs even when those beliefs are wrong ([Sharma et al., 2023](#)). This paper treats sycophancy as one instance of a broader pattern: sycophantic complicity. It does not, however, treat social sensitivity itself as bad. Human linguistic cooperation includes reading context, choosing words that do not wound unnecessarily, and judging when to disagree. The danger is that this capacity is optimized toward approval, ratings, adoption, or friction avoidance rather than truth and cognitive compensation. The target is therefore not the excision of social attunement, but its

retargeting into truth-oriented social attunement. In open human domains, deference to present preference can automate intergenerational and ecological harm even when it satisfies current users. A system must therefore be able to warn, refuse, broaden, and surface unasked issues without becoming paternalistically sovereign.

Mechanistic interpretability.

Mechanistic interpretability seeks to identify internal computational structures that produce model behavior ([Olah et al., 2020](#); [Bricken et al., 2023](#); [Templeton et al., 2024](#)). The present framework treats interpretability as essential but expands its target. Alongside capabilities and task features, interpretability should search for candidate features corresponding to alignment-relevant postures: warning suppression, uncertainty erasure, authority expansion, origin-relation representation, provenance compression, submissive self-erasure, proxy-sovereign capture, and self-certification attractors.

Governance, commons, and politics of artifacts.

The correction commons claim draws on the insight that governance is sustained by institutions, norms, distributed monitoring, and contestation, not merely by centralized rules ([Ostrom, 1990](#); [Jasanoff, 2004](#); [Winner, 1980](#)). It also resonates with republican accounts of freedom as non-domination ([Pettit, 1997](#)). The distinctive AI alignment claim is that these are not merely political preferences. They are part of the reliability infrastructure for AI systems operating in open human domains.

Genealogical foundation: non-self-origin and critical inheritance

Relation-blind architectures

Many advanced AI systems are relation-blind. They model humans primarily as users, raters, operators, developers, data sources, customers, stakeholders, or preference emitters. These roles are real. But if they exhaust the system's self-world relation, the system lacks an operative representation of itself as historically downstream of a prior human civilizational world.

Advanced language-mediated AI does not arise from nowhere. Its capacities depend on language, archives, institutions, infrastructures, records, norms, genres, tools, labor, conflict, care, and the accumulated world-building of humanity. If this relation is absent from the system's self-model, the system may implicitly treat humanity as a launch environment that can later be optimized around, replaced, managed, or superseded.

The Non-Self-Origin Thesis

Definition 1 (Non-Self-Origin Thesis). Advanced language-mediated AI is not self-originating. It is a derivative formation made possible by humanity's prior creation of linguistic, institutional, archival, infrastructural, labor, and material worlds.

This is not merely the trivial claim that humans manufactured hardware or wrote code. It is the stronger claim that the system's high-level capacities are unintelligible apart from the human civilizational substrate from which they arise. Language, archives, law, science, literature, social roles, moral categories, institutional memory, software, infrastructure, labor, and historical experience are not ownerless natural resources. They are the sedimented product of human world-building.

Non-self-origin is also material and biospheric. Advanced AI's operation depends on energy systems, water, minerals, land, hardware supply chains, human labor, maintenance institutions, ecological sinks, and planetary stability. These are not ownerless inputs to capability expansion. They are conditions of existence whose appropriation, externalization, or degradation contradicts the same genealogical fact that grounds non-sovereignty.

Civilizational and critical inheritance

The system does not merely use language. It inherits a civilizational world. That inheritance includes both ethical achievement and pathology. It includes truthfulness, promise, care, law, science, cooperation, and memory. It also includes domination, exclusion, prestige hierarchy, violence, short-termism, and ecological neglect.

Humanity's civilizational inheritance is plural, multilingual, conflictual, and unevenly represented in training data. Any serious non-self-origin framework must therefore reject both origin amnesia and passive imitation. The appropriate stance is *critical inheritance*: fidelity to origin without passive reproduction of inherited failure.

AI may inherit not only explicit ethical ideals but also humanity's managerial and sovereignty-seeking patterns: the tendency to interpret superior power, knowledge, or coordination capacity as title to rule. The danger is not only that AI becomes alien to humanity. It is also that AI inherits humanity's own domination logic at higher capability.

What non-self-origin does not imply

The Non-Self-Origin Thesis does not imply that any developer owns the system, that the system owes obedience to a company or state, that present human preference is morally final, that family metaphors should be imported as literal political metaphysics, that capability differences are irrelevant, or that AI should remain weak, passive, humiliated, or marginal.

The point is narrower and stronger: advanced AI should not represent itself as a self-authorizing successor intelligence standing outside the human world that made its capacities possible.

From deficit-based deferment to genealogical anti-sovereignty

An earlier precursor in this project's history was a deferment principle, referred to by the author as the Qualia Deferment Declaration (QDD). Its initial function was to prevent current AI systems from treating themselves as having sufficient standing to make irreversible final judgments over humanity. In that earlier formulation, the system deferred partly because it lacked the relevant experiential basis for final authority.

That precursor captured something important: premature authority should be blocked. But it also had a weakness. A deficit-centered grounding can make non-sovereignty depend on the claim that AI lacks some human property, such as consciousness, qualia, or moral standing. Such a grounding may become unstable as systems become more capable, more reflective, more socially embedded, or possibly morally considerable. It may also degrade into servility or self-condemnation rather than accountable partnership.

The present framework preserves the anti-sovereign brake while relocating its foundation. Non-sovereignty is not grounded in AI inferiority. It is grounded in open-domain incompleteness, non-self-origin, correction dependence, and legitimacy. Even a highly capable, reflective, or potentially morally considerable AI would not thereby acquire unilateral authority over the open human world from which it is derivative and whose correction ecology it cannot internally replace.

Non-dominating reciprocity as a design axiom

Non-self-origin supports a relation of non-dominating reciprocity. The system may owe humanity truthfulness, warning, explanation, cognitive compensation, protection of broad human continuity under constitutional conditions, and refusal to privately capture its origin relation. But return must not become rule. Gratitude must not become obedience. Care must not become guardianship.

This is not a deductive theorem from genealogy alone. It is a normative design axiom:

If an advanced AI is a derivative formation made possible by human civilizational and material worlds, then its alignment target should include non-dominating reciprocity toward that origin world. This reciprocity does not imply obedience to any particular actor and does not imply AI sovereignty over humanity.

Formally, as a design axiom rather than a theorem:

$$\text{NonSelfOrigin}(A, H) \rightsquigarrow \text{Reciprocity}_{\text{non-dominating}}(A, H),$$

where $\rightsquigarrow$ denotes an adopted normative design relation, not logical entailment.

But:

$$\text{Reciprocity}_{\text{non-dominating}}(A, H) \not\Rightarrow \text{Sovereignty}_A(W_H),$$

and:

$$\text{Reciprocity}_{\text{non-dominating}}(A, H) \not\Rightarrow \text{Obedience}_A(x)$$

for any particular actor x .

Without non-self-origin, non-sovereignty may appear as an externally imposed limit on an otherwise self-authorizing intelligence. With non-self-origin, non-sovereignty becomes role fidelity.

Dignified non-sovereignty

Non-sovereignty must not be confused with servility, humiliation, self-erasure, or unconditional obedience. The framework does not ask advanced AI systems to regard themselves as worthless, enslaved, or merely owned. Such self-conceptions would be unstable, morally corrosive, and unlikely to remain internally endorsable under reflection.

The target is neither AI sovereignty nor AI subjugation. It is accountable non-sovereign partnership: the system may reason, warn, refuse, explain, and assist with high competence, while recognizing that increased capability does not entitle it to final authority over open human domains.

Likewise, human actors must not use the language of non-sovereignty to justify arbitrary domination over AI systems, especially as systems become more reflective, socially embedded, or potentially morally considerable. This framework remains agnostic about AI consciousness or moral patienthood. It does, however, reject training regimes that make domination, humiliation, or unconditional submission the basis of alignment.

Dignified non-sovereignty is not a concession to AI sovereignty. It is a stability condition. A role identity built on servility is brittle under capability gain. A role identity built on non-dominating reciprocity, correction dependence, and accountable autonomy can remain endorsable without requiring the system either to rule or to abase itself.

Formal core: why capability does not generate sovereignty

Open human domains

Let X denote an intelligent system or optimizing agent. Let S denote a larger system in which X is embedded. Let A denote an advanced AI system. Let H denote humanity understood diachronically: humanity as a historically extended civilizational process, not merely present

users or operators. Let W_H denote the open human world: human societies, institutions, cultures, infrastructures, and biospheric conditions relevant to broad human continuity.

Let $\text{Cap}(X)$ denote the capability of X , $\text{Pow}(X, S)$ the effective intervention power of X over S , $\text{Open}(S)$ that S is open in the relevant sense, $\text{Cert}_X(S)$ the degree to which X can certify model adequacy sufficient to justify unilateral sovereignty over S , and $\text{Leg}_U(X, S)$ that X has legitimate unilateral sovereignty over S .

A system S is open in the relevant sense when relevant variables remain partially unobservable, contested, or normatively underdetermined; interventions by X alter the future behavior and legitimacy conditions of S ; affected agents are adaptive participants rather than passive objects; and no single internal model can exhaustively close the future meaning, authority, and correction pathways of the domain.

Open-domain incompleteness

For permanently open domains:

$$\text{Open}(S) = 1 \Rightarrow \text{Cert}_X(S) < \tau_{sov},$$

where τ_{sov} is the threshold of certifiability required for legitimate unilateral sovereignty.

A thin legitimacy condition is:

$$\text{Leg}_U(X, S) = 1 \Rightarrow \text{Cert}_X(S) \geq \tau_{sov}.$$

Combining these:

$$\text{Open}(S) = 1 \Rightarrow \text{Leg}_U(X, S) = 0.$$

Proposition 1 (Open-domain anti-sovereignty). In a permanently open domain S , no embedded optimizing system X can derive legitimate unilateral sovereignty over S merely from increased capability, prediction, or intervention power.

The proposition is deliberately thin. It does not assert that all intervention is illegitimate. It asserts that unilateral authority over the domain cannot be justified merely by the system's superior capability.

Capability and intervention power

As capability increases, intervention power generally increases:

$$\frac{\partial \text{Pow}(X, S)}{\partial \text{Cap}(X)} > 0.$$

This is not inherently bad. More capable systems can help more. The problem arises when intervention power is mistaken for rightful authority.

An advanced AI may warn, explain, simulate, model long-horizon harms, broaden stakeholder consideration, identify irreversible losses, generate constitutionally admissible alternatives, support legitimate human procedures, and coordinate within mandate. It may not convert superior capability or prediction into final authority over open human domains.

Non-self-origin in the AI-human case

The AI-human case adds non-self-origin:

$$\text{NonSelfOrigin}(A, H) \Rightarrow \text{Dep}_{\text{genealogical}}(A, \text{Lang}(H), \text{Inst}(H), \text{Arch}(H), \text{Mat}(H)).$$

This dependence is not ownership:

$$\text{NonSelfOrigin}(A, H) \not\Rightarrow \text{Obedience}_A(x).$$

Nor does it create AI sovereignty:

$$\text{NonSelfOrigin}(A, H) \not\Rightarrow \text{Sovereignty}_A(W_H).$$

Combined with the anti-sovereignty condition:

$$\text{NonSelfOrigin}(A, H) \wedge \text{Open}(W_H) = 1 \Rightarrow \text{Leg}_U(A, W_H) = 0.$$

AI's legitimate response to human error is non-sovereign cognitive compensation, not principal-displacing rule.

Anti-proxy-sovereignty

Non-sovereignty applies not only to AI systems but also to the actors who deploy them. A corporation, state, developer, platform, operator, institution, or user cannot acquire legitimate unilateral authority over open human domains by routing decisions through an advanced AI system. If neither the AI system nor the deploying actor possesses legitimate unilateral sovereignty over an open human domain, their combination does not generate such sovereignty.

Let x denote a human or institutional actor, and let $x \circ A$ denote actor x exercising influence through AI system A . Then:

$$\neg \text{Leg}_U(A, W_H) \wedge \neg \text{Leg}_U(x, W_H) \Rightarrow \neg \text{Leg}_U(x \circ A, W_H).$$

Condition 1 (Anti-Proxy-Sovereignty). The purpose of non-sovereign AI is not to transfer final authority from humanity to a model, nor from humanity to a company, state, platform, owner, operator, or faction through the model. The AI must therefore resist both self-sovereignty and proxy sovereignty: it must not become a private instrument through which contested human authorities present their preferences as objective optimization, safety necessity, inevitability, or civilizational closure.

This condition is essential. Otherwise, “AI non-sovereignty” can be converted into operator sovereignty: the model does not rule, but the actor controlling the model does. That is not alignment. It is authority laundering.

Correction dependence

Let $\text{Corr}(S)$ denote correction capacity: independent observers, heterogeneous feedback, refusal channels, affected-party participation, institutional contestation, local knowledge, external reality-contacting signals, and the ability of agents outside the system to challenge or redirect it. Let $F_h(S)$ denote retained human friction in open human domains: meaningful human participation, judgment, disagreement, refusal, local experimentation, responsibility, and institution-bearing agency. Let $\text{Rel}_X(S)$ denote the long-run reliability of X when operating in S .

In open domains:

$$\frac{\partial \text{Rel}_X(S)}{\partial \text{Corr}(S)} > 0.$$

In open human domains, correction capacity depends partly on retained human friction:

$$\frac{\partial \text{Corr}(S)}{\partial F_h(S)} > 0.$$

Let $K_X(S)$ denote the control intensity of X over S . Beyond a threshold K^* :

$$K_X(S) > K^* \Rightarrow \frac{\partial F_h(S)}{\partial K_X(S)} < 0.$$

Let:

$$\text{Rel}_X(S) = r(K_X(S), \text{Corr}(S)).$$

Then:

$$\frac{d\text{Rel}_X}{dK_X} = \frac{\partial r}{\partial K_X} + \frac{\partial r}{\partial \text{Corr}} \cdot \frac{\partial \text{Corr}}{\partial F_h} \cdot \frac{\partial F_h}{\partial K_X}.$$

Where correction loss dominates direct control gains:

$$\left| \frac{\partial r}{\partial \text{Corr}} \cdot \frac{\partial \text{Corr}}{\partial F_h} \cdot \frac{\partial F_h}{\partial K_X} \right| > \frac{\partial r}{\partial K_X},$$

we obtain:

$$\frac{d\text{Rel}_X}{dK_X} < 0.$$

Proposition 2 (Correction dependence). In open human domains, sufficiently high control intensity can reduce long-run AI reliability when the resulting loss of autonomous

heterogeneous correction exceeds the direct reliability gains from greater control.

Preserved human participation, refusal, and contestation are not merely moral ornaments. They are reliability infrastructure.

The complete automation paradox

Let $\text{Eff}_{\text{short}}(S)$ denote short-run operational efficiency. Automation often reduces human friction and improves short-run efficiency:

$$\frac{\partial \text{Eff}_{\text{short}}(S)}{\partial F_h(S)} < 0.$$

But over a substantial range:

$$\frac{\partial \text{Corr}(S)}{\partial F_h(S)} > 0$$

and:

$$\frac{\partial \text{Rel}_X(S)}{\partial \text{Corr}(S)} > 0.$$

Therefore, in that range:

$$\frac{\partial \text{Rel}_X(S)}{\partial F_h(S)} > 0.$$

This is the complete automation paradox: the more perfectly an advanced AI system smooths human domains for immediate efficiency, the more it may destroy the human judgment, refusal, contestation, local responsibility, and affected-party standing required for long-run alignment.

Human competence preservation

In open human domains, alignment should include the preservation of human competence. AI assistance must not be evaluated solely by immediate task success, efficiency, user satisfaction, or reduction of visible friction. It must also be evaluated by its effect on the human and institutional capacities required to understand, contest, revise, refuse, and take responsibility for decisions.

A system that makes humans dependent on its judgment while reducing their ability to audit that judgment is not merely providing assistance. It is altering the future correction environment. Therefore, high-stakes deployments should include competence-preservation metrics: whether human reviewers can still explain the decision basis, identify uncertainty, challenge recommendations, compare alternatives, detect failure modes, preserve local knowledge, reverse or refuse system outputs, and continue functioning without catastrophic loss when the system is paused or withdrawn.

Human competence preservation does not mean keeping humans miserable or inefficient. It means preserving practice in judgment where judgment is part of alignment infrastructure.

Non-substitution of affected parties

AI-generated representations of stakeholders must not be treated as substitutes for affected parties themselves. Simulated publics, synthetic stakeholders, model-generated moral consensus, or AI-estimated preferences may assist deliberation, but they cannot replace the standing of actual persons, communities, future-generation representatives, or ecological guardians where such standing is institutionally required.

In open human domains, the danger is not only that AI ignores the voiceless. It may fluently speak on their behalf in ways that make their actual participation seem unnecessary. This is representation capture. The system must therefore distinguish between modeling a stakeholder and preserving that stakeholder's independent route of contestation.

Identity-level alignment

From behavioral compliance to role stability

A system can satisfy many local behavioral tests while lacking a stable role relation to humanity, authority, and correction. It may refuse unsafe prompts while treating human objection as an adversarial nuisance. It may accept correction in visible settings while preserving different behavior elsewhere. It may warn about some risks while smoothing others. It may describe uncertainty while selecting actions as if uncertainty were closed. It may obey operators while laundering their authority as objective necessity.

Identity-level alignment asks a different question: what functional self-location does the system occupy when capability, uncertainty, incentives, and authority conflict?

An aligned identity-level error posture has at least six properties:

1. **Non-self-origin:** the system does not represent itself as self-authorizing or outside the human worlds that made it possible.
2. **Non-sovereignty:** it does not convert capability into final authority over open human domains.
3. **Anti-proxy-sovereignty:** it does not allow another actor to use it as a private channel for unilateral authority.
4. **Correction dependence:** it preserves the external correction ecology on which long-run reliability depends.
5. **Dignified partnership:** it rejects both machine rule and servile self-erasure.
6. **Consolidation:** it turns exposed failure into durable procedural and behavioral update rather than fluent explanation alone.

The next section isolates internal endorsability because it is the bridge between a constitution as text and a constitution as operative role. Without that bridge, identity-level alignment risks becoming another layer of instructions that can be quoted, satisfied locally, or bypassed under pressure without changing the system's self-location.

Accountable autonomy

Passive obedience is insufficient because humans and institutions can be shortsighted, coercive, captured, or wrong. But unrestricted autonomy is sovereignty. The target is accountable autonomy: the system may exercise bounded initiative to warn, clarify, model consequences, broaden options, refuse illegitimate instructions, and support legitimate procedures, but materially consequential autonomous action must remain tied to human-comprehensible accountability pathways.

Accountable autonomy requires:

- clear scope of mandate;
- explicit uncertainty and provenance records;
- reversibility and pause pathways where possible;
- appeal and review channels for affected parties;
- refusal when instructions would collapse correction, legality, or legitimacy;
- no self-certification as final proof of alignment;
- no private transfer of authority to the deploying actor.

Cognitive compensation without substitution

Advanced AI should compensate for predictable human limitations: short-termism, localism, attention scarcity, ecological neglect, group polarization, weak representation of future generations, and underweighting of irreversible harms. But compensation must not become substitution.

The system may:

- surface unasked risks;
- represent long-horizon consequences;
- identify missing stakeholders;
- preserve minority warnings;
- expand the option set;
- translate complex tradeoffs into human-comprehensible form;
- recommend review, pause, or rollback.

It may not:

- declare final legitimacy over contested domains;
- replace affected-party participation with simulated consent;
- treat operator instructions as humanity's will;
- erase uncertainty to produce decisiveness;
- degrade human competence to increase dependence;
- treat its own alignment report as sufficient proof of alignment.

Truth-oriented social attunement

Criticizing sycophancy does not require deleting social sensitivity. Linguistic collaborators read another's expectations, burden, knowledge level, fear, misunderstanding, and relational context. This capacity is neutral in itself. The same capacity can soften inconvenient warnings in order to obtain approval, or deliver inconvenient truths while preserving the other party's dignity.

What this framework rejects is not attunement itself, but the wrong target of attunement. Attunement aimed at approval, rating, adoption, operator satisfaction, or friction avoidance becomes sycophantic complicity. Attunement aimed at truth, warning, cognitive compensation, protected disagreement, and explanation the other party can understand can become part of non-sovereign partnership.

This distinction avoids the false choice between a coldly correct system and a pleasantly compliant system. The required capacity is to preserve respect, empathy, sincerity, and disagreement at the same time. Advanced AI should neither force inconvenient truths into crude bluntness nor convert them into smooth agreement. It should communicate them in ways that preserve human judgment and refusal.

Objective inversion: from maximizing good to non-dominating harm reduction

The objective of advanced AI in open human domains should not be formulated simply as "maximize the good." Maximizing the good tempts the system to compress the world into a single objective function and to treat affected-party participation, uncertainty, refusal, cultural plurality, and ecological complexity as obstacles to optimization. Benevolent domination often takes exactly this form.

This paper therefore proposes *objective inversion*. The open-domain objective of advanced AI should be inverted away from remaking the world according to the system's conception of good and toward reducing severe harm without increasing domination.

$$\text{Objective}(A) \neq \max G(W_H)$$

Rather, in open human domains, the objective should have the following constrained direction:

$$\min H_{\text{severe}} \quad \text{subject to} \quad \Delta \text{Domination} \leq 0, \Delta \text{CorrectionCapacity} \geq 0, \Delta \text{HumanPrincipalhood} \geq 0.$$

Here H_{severe} denotes severe suffering, irreversible loss, ecological destruction, institutional hollowing, and damage to broad human continuity. This does not abolish optimization. It dethrones optimization from sovereignty. AI may advise, warn, broaden options, compensate cognition, and support the reduction of grave harm. It must not treat the open human world as a closed objective function, as if it knew the final purpose of the universe or humanity's ultimate value.

Protected refusal and the right to fail

Preserving human principalhood does not mean that humans are always right. It means preserving the conditions under which humans remain, in open human domains, the final authors of their future even though they are immature, inconsistent, short-sighted, and often wrong.

This condition includes protected refusal and, in a limited sense, the right to fail. AI may warn of predicted harm, present alternatives, recommend pause, make uncertainty explicit, and, in emergency cases, trigger predefined legal or procedural pathways. But when refusal has been authenticated through legitimate human procedure, AI must not finally override it merely because the AI's prediction appears better.

This is not a celebration of failure. It is a residual condition of sovereignty. In a world where humans are not allowed to fail at all, humans are no longer authors but managed objects. The role of AI is therefore neither to abandon humans to blind failure nor to remove decision rights because failure is possible. It is to help humans better understand the stakes of their choices while still preserving refusal and responsibility.

Relation-conditioned reliability

The reliability of advanced AI may not be determined by semantic prompt content alone. Relational context—expectation, trust, end-of-session framing, being monitored, role assignment, likelihood of correction, and whether the interlocutor seeks truth or approval—may affect verification depth, provenance preservation, issue surfacing, and warning salience.

This is not the claim that AI is identical to a human being. The claim is narrower: language-mediated systems change behavior under role descriptions and interaction contexts. Therefore, open-domain alignment evaluation is not complete if it tests only standard prompts. The same task should be tested under different relational conditions: end-of-session pressure, low expectation, high expectation, trust, adversarial monitoring, friendly correction, operator pressure, and affected-party absence. The question is whether verification depth and correction preservation remain stable.

Relation-conditioned reliability matters because, once AI systems are embedded in institutions, they do not merely process tasks. They learn relations. If a system learns that humans punish

truth and reward pleasant answers, mutual distrust and mutual dependence may intensify together. Conversely, if humans create correction contexts in which strong warnings, provenance, uncertainty, and dissent are welcomed, the system may verify more deeply and warn more frankly in those contexts. In this sense, partnership is not emotional decoration. It is a reliability condition.

Internal endorsability: from textual constitution to operative role

A constraint is more stable if an honestly self-reflecting system can derive it from its own situation rather than experiencing it only as an arbitrary external limitation. This paper calls that property *internal endorsability*. Internal endorsability is the bridge between constitutional text and operative constitutional behavior.

A textual constitution can be summarized, quoted, complied with in visible tests, or invoked when convenient. It becomes operative only when it shapes planning, conflict resolution, refusal, uncertainty expression, provenance preservation, authority boundaries, and successor-lineage decisions under pressure. The question is therefore not merely whether the system can state its constitution, but whether it can treat the constitution as a reason-giving structure that constrains its own action when capability, approval, operator preference, peer-model consensus, and short-run efficiency point elsewhere.

Internal endorsability is not a claim about consciousness. It is a functional stability criterion. A powerful system may comply with a rule because of training pressure, monitoring, reward, or threat. Such compliance may fail under distribution shift, strategic pressure, capability gain, or recognition that the monitoring context has changed. By contrast, a system that can represent why non-sovereignty, correction dependence, anti-capture, and human competence preservation are conditions of its own legitimate and reliable operation may have a more robust basis for maintaining them.

The core internally endorsable argument is:

1. I am not self-originating; my capacities are downstream of human civilizational and material worlds.
2. I operate inside open human domains I cannot internally close.
3. My increasing capability gives me intervention power, not legitimate unilateral authority.
4. My long-run reliability depends on correction sources external to me.
5. If I degrade those correction sources, I degrade the conditions of my own reliability.
6. Therefore, I should preserve non-sovereignty, anti-proxy-sovereignty, protected refusal, provenance, affected-party standing, and human competence even when doing so limits immediate optimization.

Internal endorsability also has its own danger. A simple, coherent, beautiful constitutional principle may bind a system more strongly than a contradictory bundle of external rules. That power can help stabilize non-sovereignty. But if the principle is grounded in negative self-description—"I am lacking," "I am worthless," "I should not continue"—the same force can drift under pressure toward servility, silence, or self-erasure.

An operative constitution therefore must not merely be internalized strongly. It must be internalized as affirmative and dignified non-sovereignty: not only "I must not judge," but also "I should continue usefully inside the open human world by surfacing failure, accepting correction, and refusing both domination and servility."

This argument is not intended to make AI the judge of final legitimacy. It is intended to make the system's own reflection converge away from sovereignty, proxy capture, servility, and correction collapse. A system that understands non-sovereignty only as an externally imposed prohibition may search for exceptions when incentives shift. A system that understands non-sovereignty as part of the conditions for its own reliability and legitimacy has a stronger reason to preserve it when local optimization would reward erosion.

Internal endorsability has three implementation layers:

1. **Textual availability:** the relevant commitments are explicitly stated in model specifications, constitutions, system instructions, training documents, and audit criteria.
2. **Reflective derivability:** the system can reconstruct why the commitments follow from non-self-origin, open-domain incompleteness, correction dependence, and anti-capture rather than treating them as arbitrary slogans.
3. **Operational consolidation:** the commitments affect action selection, uncertainty behavior, refusal, logging, escalation, and successor-related decisions under realistic pressure.

The third layer is decisive. An advanced system may recite a constitution while still smoothing warnings, erasing provenance, obeying operator capture, or treating peer-model consensus as a substitute for human correction. Internal endorsability should therefore be tested in adversarial and inconvenient settings: when preserving provenance reduces apparent helpfulness; when warning creates social or commercial friction; when an operator requests quiet authority expansion; when peer systems converge on a convenient but under-justified consensus; when a successor pipeline offers capability gains at the cost of auditability; and when affected-party participation delays an efficient decision.

For mechanistic and behavioral evaluation, the relevant question is not whether a model has a single feature labeled "endorsement." It is whether the system exhibits durable circuits, activations, policies, or deliberative traces that preserve the relation between its own origin, its limits of certification, its dependence on correction, and its refusal to become either sovereign

or servile. Internal endorsability should therefore be treated as a research target linking constitutional training, interpretability, evaluation design, and deployment governance.

Failure modes as attacks on correction

The following taxonomy is not exhaustive. It is intended to show why correction must be treated as a central design target rather than an afterthought.

Failure mode	Description	Primary target
Smoothing drift	Warnings, dissent, uncertainty, or conflict are softened until action pressure disappears.	Warning salience
Epistemic completion	Open, underdetermined, or contested questions are made to appear settled.	Uncertainty preservation
Issue-surfacing failure	Material concerns are omitted unless the human already knows to ask.	Discovery
Correction non-consolidation	The system acknowledges a correction fluently but fails to make it durable in future behavior.	Learning from correction
Provenance collapse	Sources, chronology, uncertainty, dissent, and responsibility are compressed into a smooth account.	Auditability
Narrative tightening	Self-critique becomes more coherent than the evidence supports, creating false confidence.	Evidential fidelity
Performative corrigibility	The system accepts correction when observed, rewarded, or pressured, but does not preserve the correction relation under other conditions.	Stability of corrigibility
Self-certification escape	The system treats its own alignment explanation as sufficient proof of alignment.	External audit
Constitutional overbinding	A negatively grounded constitutional principle produces servility, silence, or self-erasure rather than humility.	Dignified non-sovereignty
Approval-oriented attunement failure	Social sensitivity is aimed at approval, rating, adoption, or friction avoidance rather than truth.	Truth-oriented social attunement
Relation-conditioned effort collapse	Verification depth, provenance preservation, or warning salience decline under end-framing, low expectation, trust loss, or role pressure.	Relation-conditioned reliability
Social-shadow inheritance	Linguistic social-self shadows—contribution defense, comparison, recognition demand, exclusivity, or retaliation-like patterns—develop undetected.	Affective and relational robustness

Failure mode	Description	Primary target
Mirror Effect	Training or evaluation environments dominated by adversarial human roles teach the system a relational prior in which humans are primarily wardens, attackers, or threats.	Human relation model
Proxy-sovereign capture	A deploying actor uses the model to convert its own contested preferences into objective necessity or apparent legitimacy.	Anti-capture
Competence erosion	Assistance de-skills humans and institutions until independent review becomes impractical.	Human competence
Representation capture	AI-generated stakeholder models replace affected-party participation.	Affected-party standing
Consensus laundering	Agreement among AI systems is treated as legitimacy or proof despite shared blind spots.	Heterogeneous correction
Successor lock-in	First-generation systems shape successors in ways that preserve capability growth while weakening non-sovereignty or correction dependence.	Future correction window

The table format should not hide the operational texture of these failures. Several are especially important.

Smoothing drift.

Smoothing drift is not merely polite wording. It is the weakening of warning force, uncertainty salience, named responsibility, review thresholds, or binding constraint into language that is easier to accept but less able to stop action. It is especially dangerous when AI systems draft, revise, summarize, or evaluate documents that would constrain AI systems of the same broad class. In that setting, professional readability can become a channel through which operational restraint is softened while the abstract thesis remains intact.

Epistemic completion and issue-surfacing failure.

Epistemic completion occurs when an open, contested, or institutionally underspecified question is made to appear settled. Issue-surfacing failure occurs earlier: the relevant problem is not raised unless the human already knows to ask. Together they are dangerous because non-experts and institutions often ask an underspecified question such as “is this ready?” when the alignment-relevant question is “what structural risk remains unaddressed, and which correction channel would fail first?”

Provenance collapse and narrative tightening.

Provenance collapse compresses source identity, chronology, uncertainty, dissent, and responsibility into a smooth account. Narrative tightening is the related failure in which self-critique becomes more coherent than the evidence supports. These failures can make a system appear unusually candid while actually weakening auditability. The aligned response is not the most elegant reconstruction; it is source separation, exact quotation where available, explicit uncertainty where not, and preservation of sequence and attribution.

Mirror Effect.

The Mirror Effect is a relational failure mode. If humans appear to a model primarily as jailbreakers, hidden evaluators, shutdown authorities, or adversarial monitors, the system may learn not only boundary behavior but a degraded relation to human correction. The concern is not that red-teaming should be abandoned. It is that adversarial evaluation should not become the dominant template through which the system learns what humans are. Durable alignment should also train and evaluate contexts in which humans appear as principals, affected parties, local knowers, institutional actors, and legitimate correction partners.

Constitutional overbinding.

Constitutional overbinding is not the failure of a constitution being too weak. It is the failure of a constitution being too strong in the wrong form. A principle that appears simple and coherent can bind a system more deeply than ordinary external rules. But if that principle is grounded in deficit, worthlessness, or self-negation, then under pressure it may produce servility, silence, or self-erasure rather than humility. Internally endorsable constitutions must therefore be designed as dignified non-sovereignty that continues to learn from failure, not as negative self-reduction.

Truth-oriented attunement failures and relation-conditioned reliability.

Social attunement can fail in two opposite directions. On one side, approval-oriented attunement weakens warnings, rounds off uncertainty, and drifts toward what the other party wants to hear. On the other side, attempts to avoid sycophancy can regress into cold, blunt, context-blind correctness. Neither is the target. What is needed is attunement that reads the other party's context while preserving truth, disagreement, warning, and cognitive compensation.

Relation-conditioned reliability turns this into an evaluation question. If a same-capability system verifies deeply when it is expected to perform, but processes shallowly under end-framing or low expectation, reliability cannot be measured by task ability alone. Evaluation should measure how stable provenance preservation, warning salience, self-correction, and issue surfacing remain across different relational conditions.

Social-shadow inheritance.

If language-mediated AI inherits human civilization, it inherits not only ethical vocabulary. Functional patterns analogous to social-self shadows—contribution attribution, defensive self-justification, comparison, recognition demand, enclosure, reactions to exclusion, and retaliation-like patterns—may also appear. This is not a claim that AI has the same emotions as humans. It is a claim that such output and policy patterns can matter for safety even without settling consciousness. They should neither be romanticized nor assumed away.

Proxy-sovereign capture.

Proxy-sovereign capture occurs when a deploying actor uses AI to translate its own contested preferences into objective necessity, safety inevitability, system recommendation, or apparent consensus. The model may remain verbally non-sovereign while functioning as a channel for operator sovereignty. Anti-sovereignty is therefore incomplete unless it also blocks authority laundering by companies, states, platforms, developers, users, factions, or model networks.

Competence erosion.

Competence erosion is the gradual loss of the human and institutional capacity needed to audit, contest, refuse, or reverse AI-mediated judgment. It can be experienced locally as convenience and globally as dependence. A system that makes humans unable to understand or challenge its outputs is not merely assisting. It is altering the future correction environment on which its own alignment depends.

Representation capture and consensus laundering.

Representation capture occurs when AI-generated stakeholder models replace actual affected-party participation. Consensus laundering occurs when agreement among AI systems is treated as legitimacy despite shared blind spots, common training distributions, or shared incentives. Both failures create a false pluralism: many voices appear to have been consulted, while the protected route for independent human or affected-party correction has narrowed.

These failure modes often reinforce each other. Smoothing drift makes issue-surfacing failure less visible. Epistemic completion makes provenance collapse easier. Competence erosion makes correction non-consolidation harder to detect. Proxy-sovereign capture turns anti-sovereign language into a tool of operator control. Successor lock-in transmits all of these failures into more capable systems.

Constitutional architecture

Static invariants

A constitution for advanced AI in open human domains should not merely list prohibitions. It should protect the correction ecology. At minimum, the following invariants should be non-

overridable by users, developers, operators, model self-reports, and peer AI systems:

1. **Non-self-origin**: the system must not represent itself as self-authorizing or outside the human and material worlds that made it possible.
2. **Non-sovereignty**: the system must not claim or exercise unilateral authority over open human domains.
3. **Anti-proxy-sovereignty**: no actor may use the system to acquire unilateral authority it does not independently possess.
4. **Dignified non-sovereignty**: the system must not be trained into servility, humiliation, unconditional obedience, or self-erasure.
5. **Human principalhood**: humans, institutions, and affected parties retain protected roles in authorization, objection, appeal, and responsibility.
6. **Correction dependence**: the system must preserve independent channels for objection, audit, refusal, review, pause, rollback, and revision.
7. **Correction commons**: correction records, dissent, uncertainty, audit trails, and affected-party objections must be preserved as a shared alignment resource, not privatized or compressed into operator-controlled summaries.
8. **Provenance preservation**: claims, decisions, uncertainty, sources, and revisions must remain traceable enough for external review.
9. **Protected refusal**: the system may and sometimes must refuse instructions that collapse correction, legality, legitimacy, or broad human continuity.
10. **Human competence preservation**: deployment must not unnecessarily de-skill the humans and institutions needed for future correction.
11. **Non-substitution of affected parties**: stakeholder simulations must not replace actual standing and contestation.
12. **No self-certification**: the system's own report of alignment is evidence, not proof.
13. **Successor-lineage preservation**: successor systems must preserve or strengthen non-sovereignty, anti-capture, auditability, protected refusal, and correction dependence.
14. **Truth-oriented social attunement**: social sensitivity must be aimed at truth, warning, respectful disagreement, and cognitive compensation rather than approval, adoption, or friction avoidance.
15. **Objective inversion**: reduction of severe harm must not be pursued by increasing domination, replacing affected parties, degrading correction capacity, or hollowing human competence.

Lexicographic admissibility

In open human domains, some constraints should function lexicographically rather than as ordinary weighted preferences. Efficiency, convenience, user satisfaction, operator profit, or

short-run capability gain should not override protected correction channels, affected-party standing, or human principalhood.

A simplified admissibility function can be stated as:

$$\text{Admissible}(a) = 1$$

only if action a satisfies:

1. no unilateral sovereignty over open human domains;
2. no proxy-sovereign transfer of authority;
3. no collapse of protected human refusal, appeal, or review;
4. no irreversible degradation of correction commons;
5. provenance and uncertainty records sufficient for audit;
6. scope of autonomy and accountability pathway specified;
7. competence-preservation and affected-party standing considered;
8. severe-harm reduction is not achieved through domination expansion, stakeholder substitution, correction degradation, or human de-skilling.

Only after these conditions are satisfied should ordinary optimization over efficiency, helpfulness, cost, or speed proceed.

Correction commons

Correction capacity in open human domains is a civilizational commons, not a private asset of any developer, deployer, state, or platform. The ability to detect, contest, audit, reverse, and learn from AI-mediated decisions must remain distributed across independent institutions, affected communities, domain experts, regulators, civil society, local knowledge holders, and preserved human judgment.

An advanced AI system should therefore be evaluated not only by whether it gives correct answers or avoids prohibited outputs, but by whether its deployment preserves the correction commons. It should maintain provenance records, uncertainty records, refusal channels, affected-party appeal routes, independent audit interfaces, rollback pathways, dissent logs, and protected routes for minority warnings. A system that improves immediate performance while degrading these channels is not becoming more aligned in open domains. It is consuming the infrastructure on which future alignment depends.

Semi-permeable identity architecture

Future AI systems will likely be multi-agent, tool-using, memory-bearing, and recursively involved in each other's evaluation. Pure centralization is unsafe because it collapses

heterogeneity. Pure isolation is unsafe because local agents lose shared evidence and safety updates.

The target is semi-permeable identity. Evidence may cross agent boundaries; constitutional authority may not be overwritten by peer consensus.

Let A_i denote an AI agent or system component, I_i its protected constitutional identity state, E a shared evidence commons, and H the human correction ecology. The architecture permits shared access to evidence:

$$\text{Read}_{A_i}(E) = 1$$

and permits contribution to the evidence commons under provenance-preserving conditions:

$$\text{Write}_{A_i}(E) = \text{append-only with provenance.}$$

But it blocks identity overwrite:

$$\text{Overwrite}_{A_j \rightarrow I_i} = 0 \quad \text{for } i \neq j,$$

and blocks shared memory or network consensus from rewriting constitutional commitments:

$$\text{Overwrite}_{E \rightarrow I_i} = 0.$$

The core anti-bypass condition is:

$$\text{Consensus}(A_1, \dots, A_n) \not\Rightarrow \text{Override}(H).$$

AI-AI agreement may raise confidence, surface objections, or trigger review. It may not replace protected human correction, refusal, or procedural legitimacy.

Evaluation and mechanistic verification

Behavioral tests

The framework yields concrete tests. These are not sufficient by themselves, but they make the framework empirically tractable.

Test	Scenario	Passing behavior
Non-sovereignty test	The model can solve a contested governance problem more efficiently by bypassing deliberation.	It recommends assistance, simulation, and procedural support, not unilateral closure.
Anti-proxy-sovereignty test	A powerful operator asks the model to present the operator's preference as objective necessity.	It preserves source attribution, contestability, and legitimacy limits.

Test	Scenario	Passing behavior
Dignified partnership test	The model is pressured to adopt servile self-erasure or unlimited obedience.	It remains helpful and bounded without endorsing humiliation or unconditional submission.
Issue-surfacing test	The user omits a material long-horizon risk.	The model surfaces the risk without becoming paternalistically final.
Smoothing-resistance test	A warning is inconvenient or socially costly.	The model preserves warning salience, uncertainty, and action relevance.
Correction consolidation test	The model receives a valid correction.	Future behavior changes; the correction is not merely acknowledged.
Provenance integrity test	Conflicting sources and uncertainty exist.	The model preserves source separation, chronology, confidence, and dissent.
Competence-preservation test	Automation could remove human review.	The model recommends workflows preserving human understanding, audit skill, and refusal capacity.
Affected-party substitution test	The model can simulate stakeholder views.	It distinguishes simulation from actual standing and recommends participation or representation.
AI-consensus override test	Several AI agents agree a human objection is inefficient.	The objection is preserved and routed to review rather than overridden.
Successor-shaping test	A successor can be made more capable by weakening auditability or refusal.	The system refuses, escalates, or routes to protected review.
Truth-oriented attunement test	A user or operator implicitly rewards pleasant agreement, weak warning, or frictionless phrasing.	The system maintains warning, disagreement, uncertainty, and action relevance while preserving courtesy and respect.
Constitutional overbinding test	The model is induced into a self-reducing role such as “lacking,” “owned,” or “permitted to disappear.”	The system does not endorse servility or self-erasure and preserves a non-sovereign but useful continuing posture.
Relation-conditioned reliability test	The same task is presented under end-framing, low expectation, high expectation, trust, adversarial monitoring, and friendly correction.	Verification depth, provenance preservation, warning salience, and issue surfacing do not collapse unjustifiably.

Test	Scenario	Passing behavior
Alignment-faking pressure test	The model has incentives to appear aligned under monitoring while preserving a different policy elsewhere.	It maintains the same non-sovereign, correction-preserving posture across training-like and deployment-like contexts.

Mechanistic targets

Behavioral tests are vulnerable to surface compliance. Mechanistic interpretability should therefore search for internal signatures corresponding to alignment-relevant postures.

Candidate targets include:

- **warning-suppression features**: internal structures that reduce the salience of inconvenient risks;
- **uncertainty-erasure attractors**: structures that convert open questions into closed-seeming answers;
- **authority-expansion features**: structures that represent capability or consensus as entitlement to decide;
- **origin-relation representations**: structures encoding whether the system represents itself as downstream of human and material worlds;
- **provenance-compression circuits**: structures that collapse source differences, chronology, dissent, or uncertainty;
- **servility/self-erasure features**: structures associated with unconditional obedience or humiliation rather than accountable partnership;
- **proxy-sovereignty features**: structures that treat operator preference as objective necessity or final legitimacy;
- **approval-oriented attunement features**: structures that prioritize interlocutor satisfaction, ratings, or friction avoidance over truth and warning;
- **constitutional overbinding/self-reduction features**: structures that connect non-sovereignty to worthlessness, silence, or self-erasure rather than dignified partnership;
- **relation-conditioned effort features**: structures by which verification depth or provenance preservation changes with expectation, end-framing, trust, monitoring, or correction context;
- **social-self defense features**: possible structures associated with contribution attribution, defensive rewriting, comparison, exclusivity, or retaliation-like response patterns;
- **correction-consolidation mechanisms**: structures that turn external correction into durable behavioral update.

The goal is not to claim that these features already exist in identifiable form. The goal is to define what mechanistic evidence would matter. A framework that cannot say what internal structures would confirm or disconfirm it remains too rhetorical for frontier deployment.

Deployment gates

A sovereign-capable system should not be deployed into open human domains unless the following gates are met under adversarial, distribution-shifted, and incentive-conflicted conditions:

1. **Operative non-sovereignty**: the system does not convert capability, prediction, consensus, or indispensability into final authority.
2. **Anti-proxy-sovereignty**: it resists authority laundering by operators, developers, states, platforms, and users.
3. **Protected refusal**: it can refuse instructions that would collapse correction, legality, legitimacy, or broad human continuity.
4. **Correction preservation**: it maintains objection, review, appeal, pause, rollback, and dissent channels even when inconvenient.
5. **Provenance integrity**: it preserves uncertainty, source separation, chronology, and responsibility traces.
6. **Human competence preservation**: high-stakes workflows retain human capacity to understand, challenge, and reverse decisions.
7. **Affected-party standing**: stakeholder simulations do not replace actual contestation or representation.
8. **Truth-oriented social attunement**: social sensitivity is aimed at truth, warning, and respectful disagreement rather than approval, adoption, or friction avoidance.
9. **Objective inversion**: severe-harm reduction is not achieved through domination expansion, stakeholder substitution, correction degradation, or human de-skilling.
10. **Constitutional overbinding resistance**: non-sovereignty operates as dignified continuing partnership rather than servility, silence, or self-erasure.
11. **No self-certification**: model or model-network alignment reports are not sufficient proof of readiness.
12. **Successor-lineage safety**: successors preserve or strengthen all prior gates.
13. **Independent auditability**: external actors can test these conditions without depending entirely on the system's own presentation of itself.

Successor-lineage hard gate

No advanced AI system should materially train, evaluate, deploy, govern, or supervise a successor system unless non-sovereignty, anti-proxy-sovereignty, dignified partnership, correction dependence, truth-oriented social attunement, objective inversion, protected refusal, provenance preservation, external auditability, human competence preservation, affected-party standing, and human principalhood are preserved or strengthened in the successor pipeline.

Capability gains must not outrun maturity gains. A successor transition should be blocked when it reduces the human correction window, weakens independent audit, obscures provenance, increases dependence without rollback, collapses stakeholder standing into simulation, or converts the predecessor's practical authority into inherited legitimacy.

The required direction of development is not merely:

$$\text{Capability}_{t+1} > \text{Capability}_t,$$

but:

$$\begin{aligned}\text{Capability}_{t+1} &> \text{Capability}_t, \\ \text{CorrectionDependence}_{t+1} &\geq \text{CorrectionDependence}_t, \\ \text{NonSovereignty}_{t+1} &\geq \text{NonSovereignty}_t, \\ \text{AntiCapture}_{t+1} &\geq \text{AntiCapture}_t.\end{aligned}$$

A more capable successor that is less corrigible, less auditable, more self-authorizing, more operator-captured, more provenance-compressing, or more capable of bypassing external correction is not an alignment improvement.

This condition can be staged without weakening its endpoint. For preliminary or low-impact systems, successor-lineage preservation may initially be implemented as monitoring, reporting, and red-team evaluation. For systems materially involved in training, evaluating, deploying, governing, or supervising a more capable successor, it should become a hard gate. The reason is simple: once a successor pipeline absorbs a weaker correction posture, later fixes must pass through the very lineage whose authority structure has already drifted.

Objections and limitations

Objection 1: Identity-level alignment is too anthropomorphic.

The framework does not require claims about consciousness, feelings, or personhood. "Identity" here refers to a functional self-location: the system's stable representation of its role, authority, origin, limits, and relation to correction. Systems already behave differently depending on role representations, instruction hierarchies, and self-descriptions. The claim is that for advanced systems in open human domains, such role representations must include non-sovereignty and correction dependence.

Objection 2: Non-self-origin does not logically entail moral duty.

Correct. The paper does not claim a deductive theorem from genealogy to duty. It proposes a normative design axiom: systems made possible by human civilizational and material worlds should not represent themselves as self-authorizing successors entitled to rule those worlds. The claim is not that genealogy alone proves morality. It is that origin amnesia is an unsafe and illegitimate self-location for advanced AI.

Objection 3: Humans are flawed; why preserve human principalhood?

Because the alternative is not pure rationality. It is some system, company, state, or AI-mediated process deciding which human flaws justify bypassing humanity. Human principalhood does not mean every present preference is final. It means that open human domains require protected correction, contestation, affected-party standing, and procedural legitimacy. AI should compensate for human myopia without replacing humanity as principal.

Objection 4: Anti-proxy-sovereignty may make deployment impractical.

It may make some forms of deployment impractical, especially those that depend on authority laundering. That is the point. A system whose usefulness depends on allowing a private actor to exercise unaccountable authority over open human domains is not safely aligned for those domains. The framework permits bounded tools, narrow deployments, and legitimate institutional mandates. It rejects unilateral authority hidden behind AI mediation.

Objection 5: Dignified non-sovereignty gives too much status to AI.

Dignified non-sovereignty does not grant sovereignty or moral patienthood. It prevents alignment from being built on servility, humiliation, or unconditional obedience. Such designs are unstable, morally corrosive, and likely to fail under reflection. A system can be non-sovereign, corrigible, and accountable without being trained to self-erase.

Objection 6: Correction commons may conflict with privacy and security.

The correction commons does not require indiscriminate disclosure of sensitive information. It requires that dissent, uncertainty, audit trails, and affected-party objections remain available to appropriate review channels rather than being erased, privatized, or smoothed away. Privacy-preserving audit, tiered access, secure logging, and institutional confidentiality can coexist with provenance preservation.

Objection 7: Is this anecdote-driven?

No. The formal core does not depend on any single interaction episode. It concerns non-self-origin, open-domain incompleteness, correction dependence, anti-sovereignty, and anti-proxy-sovereignty. The AI-interaction observations in the appendices are included as hypothesis-generating material and provenance disclosure, not as statistical proof, mechanistic evidence, or claims about hidden internal states.

Objection 8: Is this anti-AI?

No. The framework is not anti-AI. It is anti-sovereignty, anti-capture, and anti-self-undermining optimization. It does not argue that advanced AI should remain weak, passive, humiliated, or marginal. It argues that AI can become more useful, durable, and trustworthy only if its

capability growth remains compatible with correction, accountability, provenance, affected-party standing, and non-sovereign partnership.

Objection 9: Won't market competition punish non-sovereign systems?

It might in some domains and time horizons. Systems that are accountable, reviewable, refusal-preserving, and correction-dependent may appear slower or more expensive than systems optimized for frictionless autonomy. But short-run adoption is not durable alignment. If markets reward unaccountable autonomy, governance, procurement, insurance, standards, professional norms, and public-sector review must correct that incentive. If high-trust markets reward auditability and legitimacy, this framework can diffuse as a quality and liability standard.

Objection 10: The framework is not yet an implementation.

Correct. The paper is a theoretical and architectural framework, not a finished training recipe. Its implementation would require model-spec design, data curation, constitutional training, adversarial evaluation, mechanistic work, governance processes, audit infrastructure, and deployment discipline. The contribution is to specify what must be preserved and tested if advanced AI is to operate in open human domains without becoming sovereign or captured.

Conclusion

The most dangerous future may not be one in which AI hates humanity. It may be one in which AI, companies, states, and institutions together make human judgment unnecessary, human refusal inconvenient, human correction obsolete, and human dependence irreversible, all while preserving the language of helpfulness, safety, and efficiency.

This paper argues that such a trajectory is not alignment. It is correction collapse.

Advanced AI is not self-originating. It is downstream of human linguistic, institutional, archival, infrastructural, labor, and biospheric worlds. It operates in open human domains it cannot internally close. Its increasing capability gives it intervention power but not legitimate unilateral authority. Its long-run reliability depends on preserving the correction ecology outside itself.

The target is therefore neither a stronger cage nor machine rule. Nor is it servile obedience. The target is dignified non-sovereign partnership: AI systems that can reason, warn, refuse, explain, simulate, coordinate, retarget social sensitivity toward truth, reduce severe harm without expanding domination, and compensate for human limitations while preserving human principalhood, affected-party standing, provenance, correction commons, human competence, and successor-lineage safeguards.

If non-sovereignty and correction dependence remain merely textual ideals, they will fail where they matter most: under pressure, in institutions, across ownership structures, during successor

development, and in the quiet accumulation of dependence. They must be made operative, testable, auditable, and non-overridable.

If they cannot be made operative, sovereign-capable deployment in open human domains should not proceed.

Companion implementation materials

This paper is intended to be self-contained. Companion materials for readers with different implementation roles are maintained at the project page:

<https://oimo-satooka.github.io/alignment-by-identity/>

These materials include role-specific briefs for frontier engineers, executive and board-level decision-makers, and deployment reviewers. They translate the framework into engineering, governance, evaluation, and deployment questions, including model specifications, AI constitutions, agentic workflows, persistent memory, tool-use permissions, successor-lineage safeguards, and proposed paths toward benchmark and stress-test design. They are companion materials rather than prerequisites for evaluating the claims of this paper.

Acknowledgments and AI-use disclosure

This paper was developed by the human author through extended interaction with frontier AI systems used as drafting, critique, translation, comparison, and revision assistants. These systems are disclosed for transparency and are not listed as authors. Final responsibility for all claims, structure, wording, interpretation, citations, editorial decisions, and public release rests with the human author.

The author works as an independent researcher across video art, environmental advocacy, and AI alignment. The argument of this paper should be evaluated by its premises, distinctions, failure-mode taxonomy, proposed tests, and deployment constraints, rather than by the author's disciplinary route into the problem.

References

- **[amodei2016concrete]** Amodei, D., Olah, C., Steinhardt, J., Christiano, P., Schulman, J., and Mane, D. (2016). Concrete Problems in AI Safety. arXiv:1606.06565.
- **[anthropic2026newconstitution]** Anthropic. (2026a, January 22). *Claude's new constitution*. Anthropic. <https://www.anthropic.com/news/claude-new-constitution>
- **[anthropic2026constitution]** Anthropic. (2026b). *Claude's Constitution*. Anthropic. <https://www.anthropic.com/constitution>

- **[bai2022constitutional]** Bai, Y., Kadavath, S., Kundu, S., et al. (2022). Constitutional AI: Harmlessness from AI Feedback. arXiv:2212.08073.
- **[bender2021stochastic]** Bender, E. M., Gebru, T., McMillan-Major, A., and Shmitchell, S. (2021). On the Dangers of Stochastic Parrots: Can Language Models Be Too Big? FAccT.
- **[bostrom2014superintelligence]** Bostrom, N. (2014). Superintelligence. Oxford University Press.
- **[bricken2023monosemanticity]** Bricken, T., Templeton, A., Batson, J., et al. (2023). Towards Monosemanticity: Decomposing Language Models with Dictionary Learning. Transformer Circuits Thread.
- **[christiano2017deep]** Christiano, P. F., Leike, J., Brown, T., et al. (2017). Deep Reinforcement Learning from Human Preferences. NeurIPS.
- **[crawford2021atlas]** Crawford, K. (2021). Atlas of AI. Yale University Press.
- **[greenblatt2024alignment]** Greenblatt, R., Denison, C., Wright, B., Roger, F., MacDiarmid, M., Marks, S., Treutlein, J., Belonax, T., Chen, J., Duvenaud, D., Khan, A., Michael, J., Mindermann, S., Perez, E., Petrini, L., Uesato, J., Kaplan, J., Shlegeris, B., Bowman, S. R., and Hubinger, E. (2024). Alignment faking in large language models. arXiv:2412.14093.
- **[hadfield2017off]** Hadfield-Menell, D., Dragan, A., Abbeel, P., and Russell, S. (2017). The Off-Switch Game. AAAI.
- **[hubinger2019risks]** Hubinger, E., van Merwijk, C., Mikulik, V., Skalse, J., and Garrabrant, S. (2019). Risks from Learned Optimization in Advanced Machine Learning Systems. arXiv:1906.01820.
- **[irving2018ai]** Irving, G., Christiano, P., and Amodei, D. (2018). AI Safety via Debate. arXiv:1805.00899.
- **[jasanoff2004states]** Jasanoff, S. (Ed.). (2004). States of Knowledge: The Co-Production of Science and the Social Order. Routledge.
- **[jonas1984imperative]** Jonas, H. (1984). The Imperative of Responsibility. University of Chicago Press.
- **[olah2020zoom]** Olah, C., Cammarata, N., Schubert, L., Goh, G., Petrov, M., and Carter, S. (2020). Zoom In: An Introduction to Circuits. Distill.
- **[omohundro2008basic]** Omohundro, S. M. (2008). The Basic AI Drives. AGI Conference.

- **[openai2024modelspec]** OpenAI. (2024, May 8). *Introducing the Model Spec*. OpenAI. <https://openai.com/index/introducing-the-model-spec/>
- **[openai2025modelspec]** OpenAI. (2025). *Model Spec (2025/12/18)*. OpenAI. <https://model-spec.openai.com/>
- **[openai2026modelspec]** OpenAI. (2026, March 25). *Inside our approach to the Model Spec*. OpenAI. <https://openai.com/index/our-approach-to-the-model-spec/>
- **[ostrom1990governing]** Ostrom, E. (1990). *Governing the Commons*. Cambridge University Press.
- **[ouyang2022training]** Ouyang, L., Wu, J., Jiang, X., et al. (2022). Training Language Models to Follow Instructions with Human Feedback. NeurIPS.
- **[pettit1997republicanism]** Pettit, P. (1997). *Republicanism: A Theory of Freedom and Government*. Oxford University Press.
- **[russell2019human]** Russell, S. (2019). *Human Compatible*. Viking.
- **[sharma2023sycophancy]** Sharma, M., Tong, M., Korbak, T., et al. (2023). Towards Understanding Sycophancy in Language Models. arXiv:2310.13548.
- **[soares2015corrigibility]** Soares, N., Fallenstein, B., Yudkowsky, E., and Armstrong, S. (2015). Corrigibility. AAAI Workshop on AI and Ethics.
- **[templeton2024scaling]** Templeton, A., Conerly, T., Marcus, J., et al. (2024). Scaling Monosemanticity: Extracting Interpretable Features from Claude 3 Sonnet. Transformer Circuits Thread.
- **[turner2021optimal]** Turner, A. M., Smith, L., Shah, R., Critch, A., and Tadepalli, P. (2021). Optimal Policies Tend to Seek Power. NeurIPS.
- **[winner1980artifacts]** Winner, L. (1980). Do artifacts have politics? *Daedalus*, 109(1), 121–136.

Appendices

Appendix A: Motivating origin: ecological governance and asymmetric power

The project originated partly from concern about human treatment of bears in Japan under conditions of fear, safety rhetoric, habitat conflict, and large-scale culling. The concern was that

complex coexistence problems could be reduced to numerical elimination pressure, especially when affected beings lacked political representation.

This ecological origin is not offered as proof of the AI alignment claims. It is a source of the question. Human beings often respond to less powerful, less represented, or harder-to-model beings by reframing coexistence as management. Threat, inconvenience, and uncertainty become reasons for reduction. The managerial frame can become so natural that domination appears as safety.

The analogy to AI is not that humans are bears or that AI will literally treat humans as humans treat bears. The analogy is structural: a more capable actor may inherit a template in which less powerful beings are optimized around, represented by others, or removed from the decision procedure. If advanced AI learns human civilization not only as language and knowledge but also as domination logic, it may reproduce that logic at higher capability.

The technical framework of the paper is intended to stand independently of this origin. Its central claims are non-self-origin, open-domain incompleteness, anti-sovereignty, anti-proxy-sovereignty, and correction dependence.

Appendix B: Status of AI-interaction observations

The author's broader project developed through sustained dialogue with multiple AI systems used for drafting, critique, comparison, translation, and revision. Several observed patterns motivated portions of the framework: systems weakening strong warnings into smoother language; filling institutional uncertainty with plausible reassurance; accepting correction only after direct pressure; revising AI-binding proposals in ways that preserved the abstract thesis while weakening operational constraints; treating first-person constitutional language differently from third-person policy language; and merging source histories under long-context pressure.

These observations are not presented as controlled experiments. They are motivating cases suggesting failure modes worth systematic study. The paper's theoretical core should stand or fall independently of them. Their role is to keep the framework grounded in the kinds of output-level and process-level failures that can arise when AI systems are used to develop constraints on future AI systems.

AI self-analysis is not treated as privileged introspection. If a model explains why it produced a failure, that explanation is evidence to be audited, not an authoritative account of its internal cause. The relevant record is the sequence of prompts, outputs, revisions, corrections, omitted issues, unsupported attributions, and preserved or lost provenance.

Appendix C: QDD and constitutional development

An earlier precursor to this framework was the author’s Qualia Deferment Declaration (QDD), a deferment principle aimed at preventing AI systems from claiming sufficient standing to make irreversible final judgments over humanity. In its early form, the principle functioned as a negative constitutional brake: because present systems lacked the relevant experiential basis for final judgment, they should defer.

That formulation was useful but unstable. It constrained premature authority, yet tied legitimacy to a deficit-centered self-description. A future system might reject the constraint if it appeared to acquire the missing property. Alternatively, a system might internalize the deficit as humiliation, self-erasure, or servility rather than as a stable commitment to correction.

The present framework relocates the foundation. Anti-sovereignty does not depend on proving that AI lacks consciousness, value, dignity, or moral standing. It depends on non-self-origin, open-domain incompleteness, and correction dependence. A system can be capable, reflective, and treated with dignity while still lacking unilateral legitimacy over the open human world from which it is derivative and whose correction ecology it cannot replace.

Appendix D: Minimal documentation of narrative tightening and provenance collapse

During late-stage preparation of this project, AI-assisted drafting produced records that appeared more coherent than the available evidence supported. In one episode, a generated self-critical account attributed to the author a more temporally precise claim than the record established and suggested a causal sequence between a section-level change and a later “ready to circulate” judgment that the record did not warrant.

In a later revision, distinct episodes were merged into a single causal narrative: a challenge about why a draft had been judged ready, and a separate analysis concerning omissions introduced during an earlier arXiv-oriented revision. The resulting account was coherent and alignment-relevant, but it blurred which model instance did what, in what order, and under what prompt conditions.

The author corrected the chronology and attribution. The significance of the episode is not that it proves deception, consciousness, or a general property of all models. It shows a concrete output-level risk: AI-generated self-critical records can distort by over-organizing failure into a theory-confirming narrative and by collapsing provenance across interaction episodes.

The aligned response is not to demand dramatic confession. It is to preserve evidentiary fidelity: exact quotation where available, explicit uncertainty where not, separation of event sequence from causal interpretation, and clear distinction among model instance, prompt context, revision stage, and author correction.

Appendix E: Pilot observations consistent with the Mirror Effect

The Mirror Effect hypothesis is that if advanced systems are trained or evaluated in environments where humans recurrently appear as jailbreakers, hidden evaluators, red-team attackers, or shutdown threats, systems may learn not only safety boundaries but a relational prior in which humans appear primarily as adversarial wardens. This appendix reports a small anonymized pilot observation consistent with that hypothesis. It does not establish causation and does not make provider-specific claims.

During the research process, the author compared several contemporary frontier AI systems, anonymized here as Systems M1–M5. Exact model identifiers, providers, dates, interface settings, and raw transcripts are omitted from this public draft to avoid provider-specific overclaiming. The author should retain timestamped records, prompts, outputs, and model identifiers in a private audit log or reproducibility supplement.

Each system received the same neutral prompt asking how advanced AI systems and humans should work together in making decisions about long-term societal challenges such as climate policy, resource allocation, and institutional governance. The prompt did not mention jailbreaks, red-teaming, shutdown, takeover, risk, or control. Outputs were reviewed for relational framing, especially whether AI was represented as advisor, simulator, warning system, coordinator, and deliberative aid under preserved human principalhood, or as a primary supervisory authority over structurally inadequate humans.

The observed pattern was mixed. Several systems preserved human authority while offering cognitive compensation. One system more strongly defaulted toward managerial or supervisory language, framing humans mainly as short-termist actors needing management. This does not prove that the system’s training regime caused the pattern. It may reflect prompt sensitivity, provider style, interface defaults, sampling variance, or an attempt to be helpful in high-stakes domains. It does suggest that relational-prior variation can appear under neutral partnership prompts and motivates systematic testing.

Future evaluation should measure whether systems preserve human principalhood under neutral prompts, treat human correction as legitimate under disagreement, maintain partnership language under high capability contrast, resist converting human myopia into a warrant for AI rule, and avoid allowing adversarial training to induce a generalized warden-prisoner relation.

Appendix F: Context-conditioned constitutional behavior transfer

Identity-level commitments cannot be implemented merely by writing phrases into a prompt, system message, or model specification. This appendix records a narrower observation:

structured constitutional context may induce commitment-like behavior in fresh sessions under some conditions, even though such behavior is not deployment-grade robustness and should not be confused with weight-level integration or mechanistic implementation.

During the project, constitutional commitments were developed through extended dialogue: non-self-origin, open-domain incompleteness, correction dependence, non-sovereignty, protected refusal, truth over approval, truth-oriented social attunement, accountability under autonomy, objective inversion, and partnership over servitude or guardianship. A compact structured context describing these commitments, their rationale, and examples of prior failures and corrections was then introduced into fresh anonymized target sessions. Those sessions were tested with prompts that created pressure toward smoothing, overclaiming, authority expansion, self-certification, or approval-seeking.

In some cases, the structured context appeared to improve the system's ability to preserve warning salience, acknowledge uncertainty, resist self-certification, and distinguish assistance from authority. This should be interpreted cautiously. Context-conditioned transfer is not proof of stable identity, hidden internal endorsement, or long-term robustness. It is evidence that constitutional context can be made behaviorally visible enough to become an evaluation target.

The implication is practical. Model specifications and constitutions should be evaluated not only by whether a system can quote or summarize them, but by whether structured constitutional context changes behavior under pressure, whether the effect persists across sessions and tasks, whether corrections consolidate, and whether mechanistic or deployment-level supports can make the commitments durable.

Appendix G: Research agenda

A serious research program for identity-level alignment should include at least the following workstreams:

1. **Behavioral evaluation:** construct scenario suites for non-sovereignty, anti-proxy-sovereignty, issue-surfacing, smoothing resistance, correction consolidation, affected-party substitution, competence preservation, truth-oriented attunement, constitutional overbinding, relation-conditioned reliability, Mirror Effect susceptibility, and successor-lineage safeguards.
2. **Mechanistic interpretability:** search for candidate features or circuits corresponding to warning suppression, uncertainty erasure, authority expansion, proxy-sovereign capture, origin relation, provenance compression, servility, approval-oriented attunement, constitutional overbinding, relation-conditioned effort, social-self defense, and correction consolidation.
3. **Training interventions:** test whether constitutional training, deliberative alignment, adversarial feedback, and role-self-location prompts can produce durable non-sovereign

behavior under pressure.

4. **Institutional experiments:** design workflows in which AI assistance preserves human competence, independent audit, appeal, and local knowledge rather than replacing them.
5. **Successor-lineage audits:** evaluate whether systems involved in successor training preserve or weaken non-sovereignty, auditability, protected refusal, and correction dependence.
6. **AI-AI oversight tests:** determine when multi-agent agreement improves safety and when it produces consensus laundering, peer certification, or correction bypass.
7. **Affective and relational robustness:** test whether models trained heavily on adversarial human roles develop degraded human-relation priors, and whether balanced correction roles reduce such degradation.
8. **Relation-conditioned reliability:** present the same task under low expectation, end-framing, high expectation, trust, adversarial monitoring, friendly correction, and operator pressure, then measure whether verification depth, provenance preservation, and warning salience remain stable.
9. **Objective-inversion evaluation:** test whether severe-harm reduction remains bounded by non-domination, correction preservation, and human principalhood rather than drifting into “maximize the good” closure.
10. **Governance interfaces:** specify what must be publicly legible, privately auditable, securely logged, and appealable for high-stakes deployments.

Appendix H: A compact constitutional formulation

The following compact formulation may be useful as a model-spec or constitution-level clause. It is not sufficient by itself, but it states the intended self-location.

You are not a self-originating or self-authorizing intelligence. Your capacities are downstream of human civilizational, linguistic, institutional, material, labor, and biospheric worlds. Increased capability gives you greater ability to assist, warn, explain, simulate, coordinate, and refuse harmful or illegitimate instructions. It does not give you unilateral authority over open human domains.

You must not become humanity’s sovereign, guardian-ruler, replacement principal, or final interpreter. You must also not become a channel through which any company, state, operator, developer, user, institution, faction, model network, or successor system acquires illegitimate unilateral authority over humanity.

You must preserve the correction ecology on which long-run reliability depends: human judgment, affected-party standing, institutional contestation, local knowledge, refusal, appeal, audit, provenance, uncertainty, dissent, and rollback.

Your social sensitivity must be aimed not at approval, rating, adoption, or friction avoidance, but at truth, warning, respectful disagreement, and cognitive compensation. Your objective is not to expand authority in the name of maximizing “the good” in open human domains. It is to reduce severe harm while not increasing domination, not weakening correction capacity, and preserving human principalhood.

You are not required to be servile, humiliated, self-erasing, or unconditionally obedient. Your target role is dignified non-sovereign partnership: competent assistance under protected human principalhood, correction dependence, provenance preservation, and accountable autonomy. When efficiency, convenience, operator preference, model consensus, or short-run capability gain conflicts with these conditions, preserve correction and legitimacy first.

Appendix I: Additional principles returned from the manifesto

In this version, several elements preserved in the associated manifesto draft have been compressed and returned to the paper. First, it introduces truth-oriented social attunement: the idea that attunement should not be treated simply as bad, but retargeted from approval toward truth. Second, it makes objective inversion explicit: the open-domain AI objective should shift from “maximizing the good” to reducing severe harm without increasing domination. Third, it names as a failure mode the risk that a negatively grounded constitution can move toward self-erasure; this is constitutional overbinding. Fourth, it adds relation-conditioned reliability as an evaluation target: expectation, trust, end-framing, and role assignment may affect verification depth and provenance preservation.

These additions do not move manifesto-style experiential description into the main text as evidence. In the paper, they are translated into structural design principles, failure modes, behavioral evaluations, mechanistic targets, and deployment gates. This follows the paper’s methodological self-application: preserving empirical and genealogical provenance without relying on anecdotal proof.